

Table 1: Comparison on DynaHate. Prior baseline scores are reported from Vidgen et al. (2021) and Kim et al. (2025); Hate Expert and GPT-5.1 scores were obtained in local evaluation for comparison.

Model	Macro-F1	Setting
Hate Expert	0.914	Fine-tuned
RoBERTa (R4 Target; Vidgen et al., 2021)	0.759	Supervised
GPT-4o (Kim et al., 2025)	0.832	Zero-shot
GPT-4o-mini (Kim et al., 2025)	0.835	Zero-shot
GPT-5.1	0.810	Zero-shot